

Advanced concurrency

@ FProg-Spb

`alexander.vershilov@tweag.io`

August 31th, 2017



Advanced Concurrency

- "Parallel and Concurrent Programming in Haskell" S.Marlow
- Concurrency, not parallelism
- **Plan**
 - Basics
 - Exceptions
 - STM
 - Debugging (brief)

Concurrency - Threads

- `forkIO :: IO () → IO ThreadId`
forks new green thread.
- All threads are executed on HEC (Haskell Execution Context) or Capability N:M
- HECs are executed on a OS threads
- Haskell threads can be bound to OS thread.
- Communication: MVar and Exceptions

Threads - 2

Possibly unexpected things

- Program exits as soon as Thread 1 exits;
- Thread stack can't be deallocated if any ThreadId is referenced.

```
main :: IO ()
main = do
    finished <- newEmptyMVar
    forkIOFinally (...)
                (const $ putMVar finished ())
    takeMVar finished
```

- only N threads are running in parallel, where N is a number of HECs

Communication: MVar

MVar

Basic communication primitive value + lock

Shared state

`modifyMVar`

`withMVar`

`readMVar`

Properties:

- can use effects inside;
- awake only one locked thread;
- fairness;.

Communication box

`putMVar`

`takeMVar`

`readMVar`

Communication: MVar

Unexpected problems

- MVar is more lazy than one can expect; (see StrictMVar in *distributed-process*)
- Effects can escape in presence of exceptions;

Exceptions

- Asynchronous exceptions

- Ability to stop another thread;
- Cross-thread communication mechanism (bad).

- Working with exceptions

- Sending exceptions

`throwTo :: Exception e => ThreadId -> e -> IO ()`

- Masking exceptions

`mask :: ((forall a. IO a -> IO a) -> IO b) -> IO b`

- Handling exceptions

`catch :: Exception e => IO a -> (e -> IO a) -> IO a`

Masking example

```
modifyMVar :: MVar a -> (a -> IO (a,b)) -> IO b
modifyMVar m io =
  mask $ \restore -> do
    a      <- takeMVar m
    (a',b) <- restore (io a >>= evaluate)
              'onException' putMVar m a
    putMVar m a'
    return b
```


Exceptions are hard

- Unexpected interruptible operations;
- Exceptions in catch are masked but interruptibly
`'catch' (e -> someAction 'finally' someAction)`
- not possible to distinguish between sync and async exceptions;
- `SomeAsynchronousException`;
- Worker processes;
- *safe-exceptions* package.

STM

- Motivating examples
 - Exchange values in a hash-tables;
 - Composable select.
- Allow transactions;
- Restrict effects that can't be replayed;
- Allow merging of effects.

STM - primitives

- `atomically` – run transaction;
- `TVar` – basic primitive transaction var;
- `retry` – block!
- `orElse` – alternative;

```
select :: Foldable f :: Int -> f a -> IO (Maybe a)
select timeout xs = do
  d <- registerDelay timeout
  atomically $ Just <$> asum xs
    'orElse' (Nothing $> readTVar d >>= check)

checkValue :: TVar a -> (a -> Bool) -> STM a
checkValue tv f = do
  x <- readTVar tv
  check $ f x
  pure x
```

```

extract :: Int -> TVar (IntMap a) -> STM a
extract i t = do
  x <- readTVar t
  case IntMap.lookup i of
    Nothing -> retry
    Just v   -> modifyTVar t (IntMap.delete i) >> pure

tryExtract :: Int -> TVar (IntMap a) -> STM (Maybe a)
tryExtract i t = (Just <$> extract i t)
                <|> pure Nothing

extractOrFail :: Int -> TVar (IntMap a) -> STM a
extractOrFail i t = extract i t
                <|> throwSTM ValueNotFound

```

```
exchange :: Int
          -> TVar (IntMap a)
          -> TVar (IntMap a)
          -> STM ()
exchange i a b = do
  x <- extract i a
  y <- extract i b
  insert i a y
  insert i b x
```

Unexpected problems

- Termination is not guaranteed.
- All blocked threads are unblocked on change.
- No effects in transactions.
- Low performance under high contention.

To be done

Async

```
data Async a = Async ThreadID (STM (Either SomeException a))
```

- Swiss army knife for running asynchronous actions.
- Reuses best features of STM and MVars.
- If you'll try to implement good framework for running asynchronous computations and building process hierarchies then most likely your code will be very close or equal to *async*

debug

- *Event log*
- *Dejafu* package
- *RTS debug options -Ds*

What is missing

- Examples of efficient IPC communication.
- Lock-free operations.
- Speed comparison.

The End

Questions ?